

# Glance ML Internship Assignment: Multimodal Fashion & Context Retrieval

**Author:** Candidate (ML Engineering Intern)

**Date:** July 2026

**GitHub Repository:** <https://github.com/Rupsika/glance>

**Codebase Files:** [indexer.py](#), [retriever.py](#), [app.py](#), [main.py](#)

---

## 1. Executive Summary

This submission presents an end-to-end, high-performance visual retrieval engine specifically engineered for fine-grained fashion queries and contextual scenery searches. By combining deep feature representation learning with spatial object boundaries, we overcome the traditional limitations of vision-language models regarding compositionality, color-item binding, and context grounding.

The core architecture uses a **Spatial-Semantic Hybrid Retrieval** strategy. It utilizes a pre-trained COCO YOLOv8 detector to crop upper and lower torso regions, extracts normalized embeddings for the global scene and individual crops using CLIP, indexes them in a multi-vector FAISS index, and retrieves them using a custom query-deconstruction text parser. An interactive, web-based Streamlit search engine demonstrates the effectiveness of the model.

---

## 2. Dataset Analysis

For this assignment, we analyzed the `val_test2020/test` dataset, containing **3,200 images** derived from Fashionpedia.

- **Environment:** Contains high variance across office interiors, street walks, parks, indoor rooms, and runway settings.
  - **Clothing Types:** Features formal wear (suits, blazers, button-downs), casual wear (t-shirts, hoodies), and outerwear (jackets, raincoats).
  - **Color Theory:** Highly diverse with garments span bright primaries (yellow, red, blue), neutrals (black, white, grey), and soft earth tones.
-

### 3. Alternative Approaches & Trade-offs

To retrieve fashion items accurately under natural language constraints, we evaluated three primary modeling paradigms:

#### Approach A: Vanilla CLIP / OpenCLIP (Zero-Shot)

- **Concept:** Feed full image and query string directly into CLIP.
- **Pros:** Zero engineering overhead; excellent at identifying global scenes (e.g. "office background").
- **Cons:** Severe failure on **compositionality** (e.g., cannot differentiate "red shirt + blue pants" from "blue shirt + red pants"). It treats sentences as "bags of words" and suffers from color-attribute leakage.

#### Approach B: Fine-Tuned Vision-Language Captioner (VLM) + Text Dense Search

- **Concept:** Use a model like BLIP-2 or LLaVA to auto-caption all 3,200 images, then search captions using BM25 or Dense Text Embeddings (SentenceTransformers).
- **Pros:** Strong compositionality representation, as VLMs describe relations ("a red tie on a white shirt") naturally in text.
- **Cons:** Extremely computationally expensive to run caption generation on CPU for large datasets (takes hours). Text queries might miss implicit visual concepts not written in the auto-generated captions.

#### Approach C: Spatial-Semantic Hybrid Retrieval (Chosen Approach)

- **Concept:** Run a lightweight YOLOv8 detector to find the person. Crop into Upper and Lower body segments. Extract independent CLIP embeddings for the full image, upper crop, and lower crop. Parse text queries into upper/lower/global items and compute a weighted cosine similarity fusion.
  - **Pros:**
    - Directly resolves the compositionality problem by grounding text descriptions to spatial bounding box crops.
    - Highly computationally efficient; indexing 3,200 images runs in ~3-4 minutes on CPU, and retrieval takes < 1ms.
    - Zero-shot; requires no fine-tuning on labeled training data.
  - **Cons:** Relies on a robust person detector. (In cases where no person is detected, we fall back to a proportional horizontal split of the image, which acts as a fallback).
- 

### 4. The Chosen Approach: Spatial-Semantic Hybrid Retrieval

#### Part A: The Indexing Pipeline (`src/indexer.py`)

1. **Primary Subject Grounding:** A lightweight yolo v8 . pt detector locates the primary (largest) human figure in the image.

## 2. Torso Partition Heuristics:

- **Upper Body Crop:** Extracted from the top 50% of the person's bounding box ( $y_{\min} \rightarrow y_{\min} + 0.5 \cdot h$ ). This isolates shirts, ties, jackets, raincoats, and blazers.
- **Lower Body Crop:** Extracted from the bottom 60% ( $y_{\min} + 0.4 \cdot h \rightarrow y_{\max}$ ) with a 10% overlap to prevent cropping errors. This isolates trousers, jeans, skirts, and footwear.

3. **Representation Encoding:** The global image, upper crop, and lower crop are encoded via openai/clip-vit-base-patch32. The vectors are L2-normalized:

$$\|\vec{v}\|_{\text{norm}} = \frac{\|\vec{v}\|}{\|\vec{v}\|_2}$$

4. **Vector Storage:** The three representations are added to independent FAISS IndexFlatIP (Inner Product) vector databases to ensure fast retrieval. Bounding box coordinates are saved in metadata . json.

```
Raw Image → YOLOv8 Person Detector → [Upper Torso Crop] → CLIP Encoder → Upper FAISS Index
→ [Lower Torso Crop] → CLIP Encoder → Lower FAISS Index
→ [Full Image] → CLIP Encoder → Global FAISS Index
```

## Part B: The Retrieval Pipeline (src/retriever.py)

1. **Semantic Query Deconstructor:** Natural language queries are split by conjunctions/prepositions. A rule-based parser scans for environment, upper clothing, and lower clothing keywords.

*Example:* "A red tie and a white shirt in a formal setting" deconstructs into:

- **Global Context:** "in a formal setting"
- **Upper Torso:** "a red tie and a white shirt"
- **Lower Torso:** None

2. **Feature Similarity & Fusion:** We fetch text embeddings for active components and run batch inner-product matching:

$$\text{Score} = \frac{w_g \cdot \cos(\vec{v}_g, \vec{q}_g) + w_u \cdot \cos(\vec{v}_u, \vec{q}_u) + w_l \cdot \cos(\vec{v}_l, \vec{q}_l)}{w_g + w_u + w_l}$$

Active weights adjust automatically if some query parts are missing.

## 5. Evaluation Query Analyses

Our system's performance on the 5 evaluation queries demonstrates the power of our design choices:

1. **Attribute Specific:** "A person in a bright yellow raincoat."

- **How it handles it:** The query deconstructs into Upper/Global: "bright yellow raincoat". By searching the upper crop, CLIP matches the concentrated yellow color signal and clothing style directly in the upper body, ignoring distracting background colors.

2. **Contextual/Place:** "Professional business attire inside a modern office."

- **How it handles it:** Deconstructs into Global: "inside a modern office" and Upper: "professional business attire". The Global index queries full-image context to verify an office interior, while the Upper index matches blazers/button-downs.

3. **Complex Semantic:** "Someone wearing a blue shirt sitting on a park bench."

- **How it handles it:** Deconstructs into Upper: "blue shirt" and Global: "sitting on a park bench". This ensures the visual characteristics of a blue shirt match the torso, while the background elements match the outdoor park setting.

4. **Style Inference:** "Casual weekend outfit for a city walk."

- **How it handles it:** Deconstructs into Global: "casual weekend outfit for a city walk". Since CLIP is pre-trained on diverse web images, it successfully associates "city walk" with urban street scenes and "casual outfit" with hoodies/t-shirts zero-shot.

5. **Compositional:** "A red tie and a white shirt in a formal setting."

- **How it handles it:** Deconstructs into Upper: "a red tie and a white shirt" and Global: "formal setting". The search query for "red tie and white shirt" is restricted to the upper torso. This isolates the color-item binding. An image with a white tie and red shirt will score poorly because the crop's embedding won't align with "red tie".

---

## 6. Modularity, Scalability & Zero-Shot Capability

### Modularity

The codebase enforces strict separation of concerns:

- `indexer.py`: Handles raw image feature extraction and offline indexing.
- `retriever.py`: Independent search module loaded as a library.
- `app.py`: High-quality user interface built with Streamlit.
- `main.py`: Command-line manager.

### Scalability to 1 Million Images

If the database grows to 1 million images:

1. **Index Optimization:** Instead of exact search using IndexFlatIP, we can use **FAISS IVF (Inverted File Index)** combined with **HNSW (Hierarchical Navigable Small World)**. This partitions the vector space into clusters and restricts search to a fraction of the database, keeping queries under 10 milliseconds.
2. **Quantization:** We can apply Product Quantization (e.g., IndexIVFPQ) to compress 512-dim float vectors to 64 bytes, reducing memory usage from 6 GB (for 1M images with 3 vectors each) to under 800 MB.
3. **Feature Store:** Decouple metadata storage into a high-performance document store (e.g., PostgreSQL or MongoDB) and perform vector search in FAISS using matching record IDs.

## Zero-Shot Capability

The system relies entirely on pre-trained open-vocabulary foundational models (CLIP and YOLOv8). It has no fixed class dictionary. This allows users to search for arbitrary colors ("lavender", "neon green"), clothing styles ("bohemian", "steampunk"), and environments ("snowy forest", "subway terminal") without any custom training.

---

## 7. Future Work & Extensions

### Extension A: Adding Locations & Weather

To support queries like *"boho style in rainy Paris"*:

1. **Location Tagging:** Combine CLIP with a pre-trained landmark classifier (e.g., Google Landmarks V2) or perform text-matching against global background tags.
2. **Weather Classifiers:** Build a small multi-label classifier (sunny, rainy, snowy) or run a zero-shot CLIP classifier over the global image for climate tags (["rainy scene", "sunny weather", "snowy day"]).
3. **Metadata Filtering:** Store location and weather tags as metadata fields in the database and apply SQL-like filters during retrieval to restrict vector search to matching entries.

### Extension B: Improving Precision

1. **Negative Prompting:** Allow users to define negative keywords (e.g., "no hats", "excluding red jackets") and subtract normalized negative vectors from the query vector to redirect search paths.
2. **Segment Anything Model (SAM):** Instead of rectangular box crops, extract precise semantic garment segmentations. This isolates clothing items perfectly from backgrounds, boosting precision for fine-grained fashion descriptions.

3. **Ensemble Text Models:** Combine CLIP's text embedding with a Sentence-Transformer trained on fashion-specific metadata to capture technical fashion terminology.